

CPSC 340 Assignment 3 (due Monday, Feb 10 at 11:55pm)

Instructions

Rubric: {mechanics:5}

IMPORTANT!!! Before proceeding, please carefully read the general homework instructions at <https://www.cs.ubc.ca/~fwood/CS340/homework/>. The above 5 points are for following the submission instructions. You can ignore the words “mechanics”, “reasoning”, etc.

We use [blue](#) to highlight the deliverables that you must answer/do/submit with the assignment.

1 Finding Similar Items

For this question we’ll use the Amazon product data set¹ from <http://jmcauley.ucsd.edu/data/amazon/>. We will focus on the “Patio, Lawn, and Garden” section. You should start by downloading the ratings at <https://stanford.io/2Q7QTVu> and place the file in your `data` directory with the original filename. Once you do that, running `python main.py -q 1` should do the following steps:

- Load the raw ratings data set into a Pandas dataframe.
- Construct the user-product matrix as a sparse matrix (to be precise, a `scipy.sparse.csr_matrix`).
- Create bi-directional mappings from the user ID (e.g. “A2VNYWOPJ13AFP”) to the integer index into the rows of `X`.
- Create bi-directional mappings from the item ID (e.g. “0981850006”) to the integer index into the columns of `X`.

1.1 Exploratory data analysis

1.1.1 Most popular item

Rubric: {code:1}

Find the item with the most total stars. [Submit the product name and the number of stars.](#)

Note: once you find the ID of the item, you can look up the name by going to the url https://www.amazon.com/dp/ITEM_ID, where `ITEM_ID` is the ID of the item. For example, the URL for item ID “B00CFM0P7Y” is <https://www.amazon.com/dp/B00CFM0P7Y>.

Answer: “Classic Accessories Veranda Grill Cover, X-large” 4.4 Star

1.1.2 User with most reviews

Rubric: {code:1}

[Find the user who has rated the most items, and the number of items they rated.](#)

Answer: A100wO06OQR8BQ, rated 161 items


¹The author of the data set has asked for the following citations: (1) Ups and downs: Modeling the visual evolution of fashion trends with one-class collaborative filtering. R. He, J. McAuley. WWW, 2016, and (2) Image-based recommendations on styles and substitutes. J. McAuley, C. Targett, J. Shi, A. van den Hengel. SIGIR, 2015.

1.1.3 Histograms


Rubric: {code:2}

Make the following histograms:


1. The number of ratings per user



2. The number of ratings per item



3. The ratings themselves



Note: for the first two, use `plt.yscale('log', nonposy='clip')` to put the histograms on a log-scale. Also, you can use `X.getnnz` to get the total number of nonzero elements along a specific axis.

1.2 Finding similar items with nearest neighbours

Rubric: {code:6}

We'll use scikit-learn's `neighbors.NearestNeighbors` object to find the items most similar to the example item above, namely the Brass Grill Brush 18 Inch Heavy Duty and Extra Strong, Solid Oak Handle, at URL <https://www.amazon.com/dp/B00CFMOP7Y>.

Find the 5 most similar items to the Grill Brush using the following metrics (Make sure to include the code you have written in your pdf GradeScope submission.):

1. Euclidean distance (the `NearestNeighbors` default)
 Answer: (indices) 103866 103865 98897 72226 102810
 Answer: (items) B00IJB5MCS B00IJB4MLA B00EXE4O42 B00743MZCM B00HVXQY9A
 Code:

```
model_1=NearestNeighbors(n_neighbors=6, metric='euclidean')
model_1.fit(X.T)
euc_closest = model_1.kneighbors(grill_brush_vec.T,return_distance=False)
print("closest noramlized euclidean distances are: ", euc_closest)
```
2. Normalized Euclidean distance (you'll need to do the normalization)
 Answer: (indices) 103866 103867 103865 98068 98066
 Answer: (items) B00IJB5MCS B00IJB8F3G B00IJB4MLA B00EF45AHU B00EF3YF0Y
 Code:

```
model_2=NearestNeighbors(n_neighbors=6, metric='euclidean')
X_normalized = normalize(X.T)
model_2.fit(X.T)
euc_closest = model_2.kneighbors(grill_brush_vec.T,return_distance=False)
print("closest euclidean distances are: ", euc_closest)
```

3. Cosine similarity (by setting `metric='cosine'`)
 Answer: 103866 103867 103865 98068 98066
 Answer: (items) B00IJB5MCS B00IJB8F3G B00IJB4MLA B00EF45AHU B00EF3YF0Y
 Code:

```
model_3 = NearestNeighbors(n_neighbors=6, metric='cosine')
model_3.fit(X.T)
cosine_closest = model_3.kneighbors(grill_brush_vec.T, return_distance=False)
print("closest cosine are: ", cosine_closest)
```

Some notes/hints...

- If you run `python main.py -q 1.2`, it will grab the row of `X` associated with the grill brush. The mappers take care of going back and forth between the IDs (like “B00CFM0P7Y”) and the indices of the sparse array (0, 1, 2, ...).
- Keep in mind that scikit-learn’s `NearestNeighbors` is for taking neighbors across rows, but here we’re working across columns.
- Keep in mind that scikit-learn’s `NearestNeighbors` will include the query item itself as one of the nearest neighbours if the query item is in the “training set”.
- Normalizing the columns of a matrix would usually be reasonable to implement, but because `X` is stored as a sparse matrix it’s a bit more of a mess. Therefore, use `sklearn.preprocessing.normalize` to help you with the normalization in part 2.

Did normalized Euclidean distance and cosine similarity yields the same similar items, as expected?

Answer: Yes, they yielded the same exact similar items.

1.3 Total popularity

Rubric: {reasoning:2}

For both Euclidean distance and cosine similarity, find the number of reviews for each of the 5 recommended items and report it. Do the results make sense given what we discussed in class about Euclidean distance vs. cosine similarity and popular items?

Note: in `main.py` you are welcome to combine this code with your code from the previous part, so that you don’t have to copy/paste all that code in another section of `main.py`.

Euclidean distance: 55, 45, 1, 1, 1

Cosine similarity: 55, 91, 45, 66, 110

This makes sense because cosine similarity is better at returning more “popular” items. Three of the items given by euclidean distance only have 1 rating.

```
{X_sums = X_binary.getnnz(axis=0)
print("Number of reviews for Euclidean distance:")
print(X_sums[:, [103866]])
print(X_sums[:, [103865]])
print(X_sums[:, [98897]])
print(X_sums[:, [72226]])
print(X_sums[:, [102810]])
print("Number of reviews for cosine similarity:")
print(X_sums[:, [103866]])
print(X_sums[:, [103867]])
print(X_sums[:, [103865]])
print(X_sums[:, [98068]])
print(X_sums[:, [98066]])
```

2 Matrix Notation and Minimizing Quadratics

2.1 Converting to Matrix/Vector/Norm Notation

Rubric: {reasoning:3}

Using our standard supervised learning notation (X, y, w) express the following functions in terms of vectors, matrices, and norms (there should be no summations or maximums).

1. $\max_{i \in \{1, 2, \dots, n\}} |w^T x_i - y_i|$. Answer: $\|Xw - y\|_\infty$
2. $\sum_{i=1}^n v_i (w^T x_i - y_i)^2 + \frac{\lambda}{2} \sum_{j=1}^d w_j^2$. Answer: $(Xw - y)^T V (Xw - y) + \frac{\lambda}{2} \|w\|^2$
3. $(\sum_{i=1}^n |w^T x_i - y_i|)^2 + \frac{1}{2} \sum_{j=1}^d \lambda_j |w_j|$. Answer: $\|Xw - y\|_1^2 + \frac{1}{2} \|\Lambda w\|_1$

Note that in part 2 we give a *weight* v_i to each training example and the value λ is a non-negative scalar, whereas in part 3 we are regularizing the parameters with different weights λ_j . You can use V to denote a diagonal matrix that has the values v_i along the diagonal, and Λ as a diagonal matrix that has the λ_j values along the diagonal. You can assume that all the v_i and λ_i values are non-negative.

2.2 Minimizing Quadratic Functions as Linear Systems

Rubric: {reasoning:3}

Write finding a minimizer w of the functions below as a system of linear equations (using vector/matrix notation and simplifying as much as possible). Note that all the functions below are convex so finding a w with $\nabla f(w) = 0$ is sufficient to minimize the functions (but show your work in getting to this point).

1. $f(w) = \frac{1}{2} \|w - v\|^2$ (projection of v onto real space).

Answer: $f(w) = \sum_{j=1}^d (w_j - v_j)^2 = \frac{1}{2} ((w_1 - v_1)^2 + (w_2 - v_2)^2 + \dots + (w_d - v_d)^2)$

$$\frac{\delta f}{\delta w_1} = w_1 - v_1; \text{ in general } \frac{\delta f}{\delta w_j} = w_j - v_j, \text{ and } \nabla f = \begin{bmatrix} w_1 - v_1 \\ w_2 - v_2 \\ \vdots \\ w_d - v_d \end{bmatrix} = w - v.$$

Therefore for $\nabla f = 0$ choose $w = v$

2. $f(w) = \frac{1}{2} \|Xw - y\|^2 + \frac{1}{2} w^T \Lambda w$ (least squares with weighted regularization).

Answer: Xw is the n by 1 vector \hat{y} and $\frac{1}{2} w^T \Lambda w$ is $\frac{1}{2} \sum_{j=1}^d \Lambda_{jj} w_j^2$ so we have

$$f(w) = \frac{1}{2} \|\hat{y} - y\|^2 + \frac{1}{2} \sum_{j=1}^d \Lambda_{jj} w_j^2 = \frac{1}{2} \sum_{i=1}^n (w^T x_i - y_i)^2 + \frac{1}{2} \sum_{j=1}^d \Lambda_{jj} w_j^2$$

Note that $\sum_{i=1}^n (w^T x_i - y_i)^2 = (w^T x_1 - y_1)^2 + (w^T x_2 - y_2)^2 + \dots + (w^T x_n - y_n)^2$ so

$$f(w) = \frac{1}{2} ((w^T x_1 - y_1)^2 + (w^T x_2 - y_2)^2 + \dots + (w^T x_n - y_n)^2) + \frac{1}{2} (\Lambda_{11} w_1^2 + \Lambda_{22} w_2^2 + \dots + \Lambda_{dd} w_d^2)$$

$$f(w) = \frac{1}{2} ((w_1 x_{11} + \dots + w_d x_{1d} - y_1)^2 + (w_1 x_{21} + \dots + w_d x_{2d} - y_2)^2 + \dots + (w_1 x_{n1} + \dots + w_d x_{nd} - y_n)^2) + \frac{1}{2} (\Lambda_{11} w_1^2 + \Lambda_{22} w_2^2 + \dots + \Lambda_{dd} w_d^2)$$

$$\frac{\delta f}{\delta w_1} = (w_1 x_{11} + \dots + w_d x_{1d} - y_1) x_{11} + (w_1 x_{21} + \dots + w_d x_{2d} - y_2) x_{21} + \dots + (w_1 x_{n1} + \dots + w_d x_{nd} - y_n) x_{n1} + \Lambda_{11} w_1,$$

or

$$\frac{\delta f}{\delta w_1} = \sum_{i=1}^n w^T (x_i - y_i) x_{i1} + \Lambda_{11} w_1; \text{ in general } \frac{\delta f}{\delta w_j} = \sum_{i=1}^n w^T (x_i - y_i) x_{ij} + \Lambda_{jj} w_j, \text{ and}$$

$$\nabla f = \begin{bmatrix} \sum_{i=1}^n (w^T x_i - y_i) x_{i1} + \Lambda_{11} w_1 \\ \vdots \\ \sum_{i=1}^n (w^T x_i - y_i) x_{id} + \Lambda_{dd} w_d \end{bmatrix} = X^T (Xw - y) + \Lambda w.$$

Therefore for $\nabla f = 0$ choose $w = \frac{X^T y}{X^T X + \Lambda}$

3. $f(w) = \frac{1}{2} \sum_{i=1}^n v_i (w^T x_i - y_i)^2 + \frac{\lambda}{2} \|w - w^0\|^2$ (weighted least squares shrunk towards non-zero w^0).

$$\text{Answer: } f(w) = \frac{1}{2} [v_1 (\sum_{j=1}^d w_j x_{1j} - y_1)^2 + \dots + v_n (\sum_{j=1}^d w_j x_{nj} - y_n)^2] + \frac{\lambda}{2} \sum_{j=1}^d (w_j - w_j^0)^2$$

$$f(w) = \frac{1}{2} [v_1 (\sum_{j=1}^d w_j x_{1j} - y_1)^2 + \dots + v_n (\sum_{j=1}^d w_j x_{nj} - y_n)^2] + \frac{\lambda}{2} \sum_{j=1}^d (w_j^2 - 2w_j w_j^0 + w_j^{02})$$

$$\frac{\delta f}{\delta w_1} = \sum_{i=1}^n v_i (w^T x_i - y_i) x_{i1} + \lambda (w_1 - w_1^0); \text{ in general } \frac{\delta f}{\delta w_j} = \sum_{i=1}^n v_i (w^T x_i - y_i) x_{ij} + \lambda (w_j - w_j^0)$$

$$\nabla f = \begin{bmatrix} \sum_{i=1}^n v_i (w^T x_i - y_i) x_{i1} + \lambda (w_1 - w_1^0) \\ \vdots \\ \sum_{i=1}^n v_i (w^T x_i - y_i) x_{id} + \lambda (w_d - w_d^0) \end{bmatrix} = X^T V (Xw - y) + \lambda (w - w^0).$$

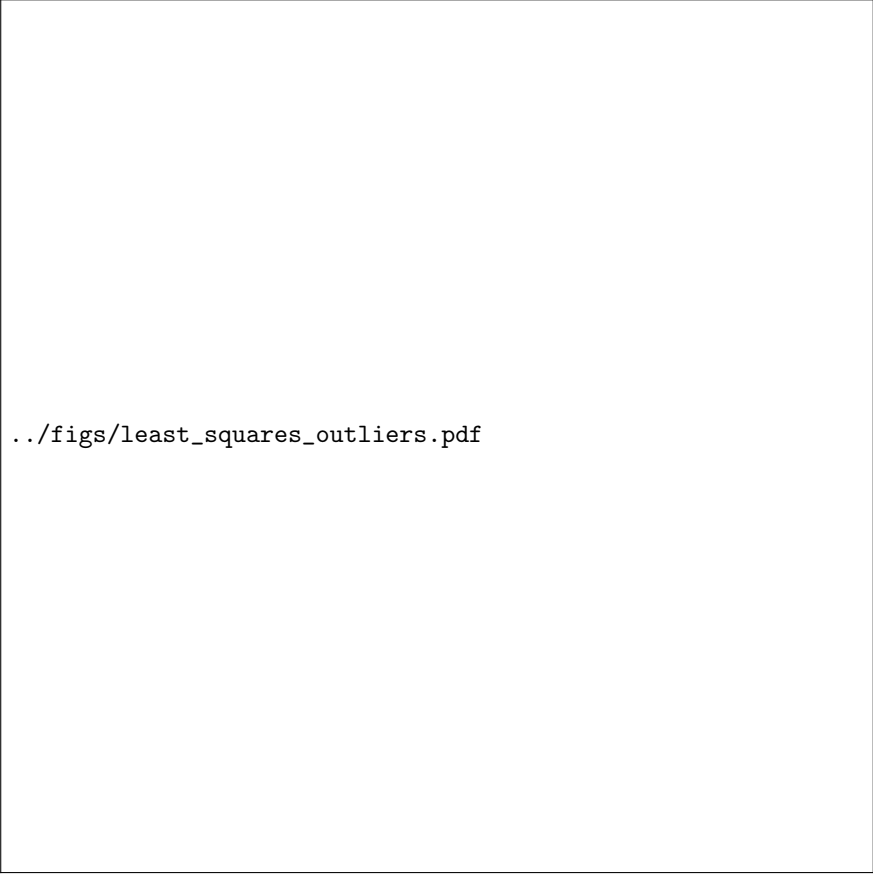
Therefore, for $\nabla f = 0$ choose $w = \frac{X^T V y - \lambda w^0}{X^T V X + \lambda}$

Above we assume that v and w^0 are d by 1 vectors (in part 3 v is a vector of length n by 1), and Λ is a d by d diagonal matrix (with positive entries along the diagonal). You can use V as a diagonal matrix containing the v_i values along the diagonal.

Hint: Once you convert to vector/matrix notation, you can use the results from class to quickly compute these quantities term-wise. As a sanity check for your derivation, make sure that your results have the right dimensions. **As a sanity check, make that the dimensions match for all quantities/operations: in order to make the dimensions match for some parts you may need to introduce an identity matrix.** For example, $X^T X w + \lambda w$ can be re-written as $(X^T X + \lambda I)w$.

3 Robust Regression and Gradient Descent

If you run `python main.py -q 3`, it will load a one-dimensional regression dataset that has a non-trivial number of ‘outlier’ data points. These points do not fit the general trend of the rest of the data, and pull the least squares model away from the main downward trend that most data points exhibit:



../figs/least_squares_outliers.pdf

Note: we are fitting the regression without an intercept here, just for simplicity of the homework question. In reality one would rarely do this. But here it’s OK because the “true” line passes through the origin (by design). In Q4.1 we’ll address this explicitly.

3.1 Weighted Least Squares in One Dimension

Rubric: {code:3}

One of the most common variations on least squares is *weighted* least squares. In this formulation, we have a weight v_i for every training example. To fit the model, we minimize the weighted squared error,


$$f(w) = \frac{1}{2} \sum_{i=1}^n v_i (w^T x_i - y_i)^2.$$

In this formulation, the model focuses on making the error small for examples i where v_i is high. Similarly, if v_i is low then the model allows a larger error. Note: these weights v_i (one per training example) are completely different from the model parameters w_j (one per feature), which, confusingly, we sometimes also call “weights”.

Complete the model class, `WeightedLeastSquares`, that implements this model (note that Q2.2.3 asks you to show how a few similar formulation can be solved as a linear system). Apply this model to the data containing outliers, setting $v = 1$ for the first 400 data points and $v = 0.1$ for the last 100 data points (which are the outliers). [Hand in your code and the updated plot](#)

Answer: Code:

```
def fit(self,X,y,z):
    X = z@X
    self.w = solve(X.T@X,X.T@y)
```



3.2 Smooth Approximation to the L1-Norm

Rubric: {reasoning:3}

Unfortunately, we typically do not know the identities of the outliers. In situations where we suspect that there are outliers, but we do not know which examples are outliers, it makes sense to use a loss function that is more robust to outliers. In class, we discussed using the sum of absolute values objective,

$$f(w) = \sum_{i=1}^n |w^T x_i - y_i|.$$

This is less sensitive to outliers than least squares, but it is non-differentiable and harder to optimize. Nevertheless, there are various smooth approximations to the absolute value function that are easy to optimize. One possible approximation is to use the log-sum-exp approximation of the max function²:

$$|r| = \max\{r, -r\} \approx \log(\exp(r) + \exp(-r)).$$

²Other possibilities are the Huber loss, or $|r| \approx \sqrt{r^2 + \epsilon}$ for some small ϵ .

Using this approximation, we obtain an objective of the form

$$f(w) = \sum_{i=1}^n \log(\exp(w^T x_i - y_i) + \exp(y_i - w^T x_i)).$$

which is smooth but less sensitive to outliers than the squared error. Derive the gradient ∇f of this function with respect to w . You should show your work but you do not have to express the final result in matrix notation.

Answer: We can rewrite $f(w)$ as:

$$\log(e^{w^T x_1 - y_1} + e^{y_1 - w^T x_1}) + \log(e^{w^T x_2 - y_2} + e^{y_2 - w^T x_2}) + \dots + \log(e^{w^T x_n - y_n} + e^{y_n - w^T x_n}).$$

Note that $w^T x_i = \sum_{j=1}^d w_j x_{ij} = w_1 x_{i1} + w_2 x_{i2} + \dots + w_d x_{id}$. Therefore for example if we take the derivative

of $f(w)$ with respect to w_1 , we get

$$\frac{\delta f}{\delta w_1} = \frac{x_{11}e^{w^T x_1 - y_1} + (-x_{11})e^{y_1 - w^T x_1}}{e^{w^T x_1 - y_1} + e^{y_1 - w^T x_1}} + \dots + \frac{x_{n1}e^{w^T x_n - y_n} + (-x_{n1})e^{y_n - w^T x_n}}{e^{w^T x_n - y_n} + e^{y_n - w^T x_n}} = \sum_{i=1}^n x_{i1} \frac{e^{w^T x_i - y_i} - e^{y_i - w^T x_i}}{e^{w^T x_i - y_i} + e^{y_i - w^T x_i}}.$$

In general $\frac{\delta f}{\delta w_j} = \sum_{i=1}^n x_{ij} \frac{e^{w^T x_i - y_i} - e^{y_i - w^T x_i}}{e^{w^T x_i - y_i} + e^{y_i - w^T x_i}}$ for some j

$$\text{Therefore the gradient } \nabla f = \begin{bmatrix} \sum_{i=1}^n x_{i1} \frac{e^{w^T x_i - y_i} - e^{y_i - w^T x_i}}{e^{w^T x_i - y_i} + e^{y_i - w^T x_i}} \\ \vdots \\ \sum_{i=1}^n x_{id} \frac{e^{w^T x_i - y_i} - e^{y_i - w^T x_i}}{e^{w^T x_i - y_i} + e^{y_i - w^T x_i}} \end{bmatrix}$$

3.3 Robust Regression

Rubric: {code:3}

The class `LinearModelGradient` is the same as `LeastSquares`, except that it fits the least squares model using a gradient descent method. If you run `python main.py -q 3.3` you'll see it produces the same fit as we obtained using the normal equations.

The typical input to a gradient method is a function that, given w , returns $f(w)$ and $\nabla f(w)$. See `funObj` in `LinearModelGradient` for an example. Note that the `fit` function of `LinearModelGradient` also has a numerical check that the gradient code is approximately correct, since implementing gradients is often error-prone.³

An advantage of gradient-based strategies is that they are able to solve problems that do not have closed-form solutions, such as the formulation from the previous section. The class `LinearModelGradient` has most of the implementation of a gradient-based strategy for fitting the robust regression model under the log-sum-exp approximation. The only part missing is the function and gradient calculation inside the `funObj` code. Modify `funObj` to implement the objective function and gradient based on the smooth approximation to the absolute value function (from the previous section). Hand in your code, as well as the plot obtained using this robust regression approach.

Answer: Code:


```
def funObj(self,w,X,y):
    n, d = X.shape
    r = X@w - y
    f = np.sum(np.log((np.exp(r)+np.exp(-r))))
    opt = np.zeros(n)
```

³Sometimes the numerical gradient checker itself can be wrong. See CPSC 303 for a lot more on numerical differentiation.

```

for i in range(n):
    opt[i] = (np.exp(w.T@X[i]-y[i])-np.exp(y[i]-w.T@X[i]))/(np.exp(w.T@X[i]-y[i])+np.exp(y[i]-w.T@X[i]))
    g = X.T@opt
return (f,g)

```



4 Linear Regression and Nonlinear Bases

In class we discussed fitting a linear regression model by minimizing the squared error. In this question, you will start with a data set where least squares performs poorly. You will then explore how adding a bias variable and using nonlinear (polynomial) bases can drastically improve the performance. You will also explore how the complexity of a basis affects both the training error and the test error.

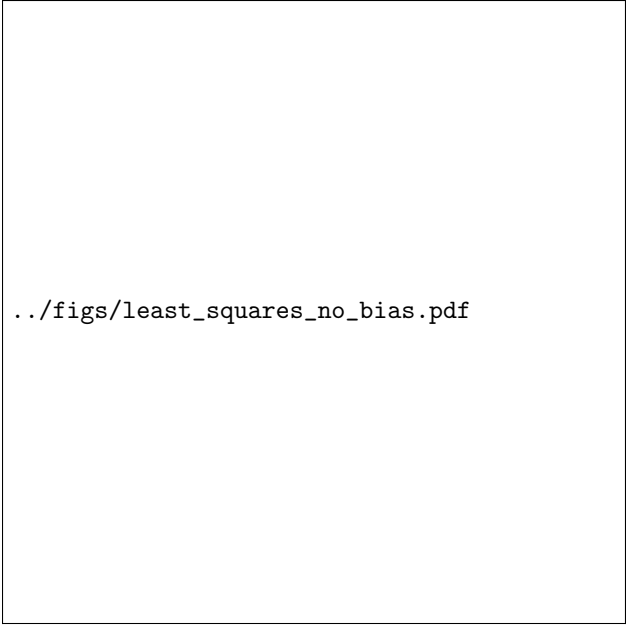
4.1 Adding a Bias Variable

Rubric: {code:3}

If you run `python main.py -q 4`, it will:

1. Load a one-dimensional regression dataset.
2. Fit a least-squares linear regression model.
3. Report the training error.
4. Report the test error (on a dataset not used for training).
5. Draw a figure showing the training data and what the linear model looks like.

Unfortunately, this is an awful model of the data. The average squared training error on the data set is over 28000 (as is the test error), and the figure produced by the demo confirms that the predictions are usually nowhere near the training data:



../figs/least_squares_no_bias.pdf

The y -intercept of this data is clearly not zero (it looks like it's closer to 200), so we should expect to improve performance by adding a *bias* (a.k.a. intercept) variable, so that our model is

$$y_i = w^T x_i + w_0.$$

instead of

$$y_i = w^T x_i.$$

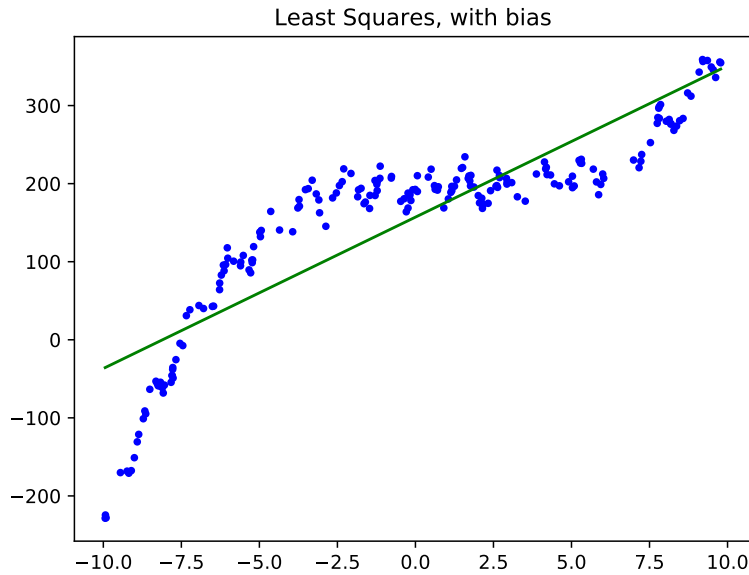
In file `linear_model.py`, complete the class, `LeastSquaresBias`, that has the same input/model/predict format as the `LeastSquares` class, but that adds a *bias* variable (also called an intercept) w_0 (also called β in lecture). Hand in your new class, the updated plot, and the updated training/test error.

Hint: recall that adding a bias w_0 is equivalent to adding a column of ones to the matrix X . Don't forget that you need to do the same transformation in the `predict` function.

Answer: Training Error: 3551.3, Test Error: 3393.9

Code:

```
class LeastSquaresBias:
    def fit(self, X, y):
        N, D = X.shape
        ones = np.ones(N)
        X = np.column_stack((ones, X))
        self.w = solve(X.T @ X, X.T @ y)
    def predict(self, X):
        N, D = X.shape
        ones = np.ones(N)
        X = np.column_stack((ones, X))
        return X @ self.w
```



4.2 Polynomial Basis

Rubric: {code:4}

Adding a bias variable improves the prediction substantially, but the model is still problematic because the target seems to be a *non-linear* function of the input. Complete `LeastSquarePoly` class, that takes a data vector x (i.e., assuming we only have one feature) and the polynomial order p . The function should perform a least squares fit based on a matrix Z where each of its rows contains the values $(x_i)^j$ for $j = 0$ up to p . E.g., `LeastSquaresPoly.fit(x,y)` with $p = 3$ should form the matrix

$$Z = \begin{bmatrix} 1 & x_1 & (x_1)^2 & (x_1)^3 \\ 1 & x_2 & (x_2)^2 & (x_2)^3 \\ \vdots & & & \\ 1 & x_n & (x_n)^2 & (x_n)^3 \end{bmatrix},$$

and fit a least squares model based on it. [Hand in the new class](#), and report the training and test error for $p = 0$ through $p = 10$. Explain the effect of p on the training error and on the test error.

Note: you should write the code yourself; don't use a library like sklearn's `PolynomialFeatures`.

Both training error and test error decrease dramatically from $p=0$ through $p=3$. As p increases further to 10, training error decreases slightly from 252.0 to 235.1. Test error decreases to 239.5 at $p=5$ and then increases to 256.3 at $p=10$.

$p=0$: Training error = 15480.5, Test error = 14390.8
 $p=1$: Training error = 3551.3, Test error = 3393.9
 $p=2$: Training error = 2168.0, Test error = 2480.7
 $p=3$: Training error = 252.0, Test error = 242.8
 $p=4$: Training error = 251.5, Test error = 242.1
 $p=5$: Training error = 251.1, Test error = 239.5
 $p=6$: Training error = 248.6, Test error = 246.0
 $p=7$: Training error = 247.0, Test error = 242.9

p=8: Training error = 241.3, Test error = 246.0
 p=9: Training error = 235.8, Test error = 259.3
 p=10: Training error = 235.1, Test error = 256.3

```
class LeastSquaresPoly:
    def __init__(self, p):
        self.leastSquares = LeastSquares()
        self.p = p
    def fit(self, X, y):
        Xi = self.__polyBasis(X)
        self.w = solve(Xi.T@Xi, Xi.T@y)
    def predict(self, X):
        Xi = self.__polyBasis(X)
        return Xi@self.w
    def __polyBasis(self, X):
        N, D = X.shape
        Z = np.ones((N, self.p + 1))
        for i in range(N):
            for p in range(Z.shape[1]):
                Z[i][p] = X[i]**p
        return Z
```

5 Very-Short Answer Questions

Rubric: {reasoning:7}

1. Suppose that a training example is global outlier, meaning it is really far from all other data points. How is the cluster assignment of this example by k -means? And how is it set by density-based clustering?
 Answer: By k -means, the outlier will be "forced" into the nearest cluster. By density-based clustering, the outlier will remain unassigned.
2. Why do we need random restarts for k -means but not for density-based clustering?
 Answer: For k -means, the first set of points where the algorithm starts is chosen randomly. It may be that the solution found is not the most optimal (i.e. a true "cluster" is split in two). Density based clustering always finds the most optimal clustering.
3. Can hierarchical clustering find non-convex clusters?
 Answer: Yes. Hierarchical clustering is also density-based. All density-based clustering techniques can find non-convex connected clusters.
4. For model-based outlier detection, list an example method and problem with identifying outliers using this method.
 Answer: We can use Z-score to check for outliers. For example, we can come up with a threshold of $|z|$ that categorizes the observation as an outlier. However, this model could be flawed because coming up with the model itself (mean and variance) is already sensitive to outliers.
5. For graphical-based outlier detection, list an example method and problem with identifying outliers using this method.
 Answer: Plot the data in a 2D scatterplot. Look at the plot, and decide if any points are outliers based on qualitative analysis. The problem with this is that only 2 (maybe 3 in a 3D plot) variables can be analyzed at a time.
6. For supervised outlier detection, list an example method and problem with identifying outliers using this method.

Answer: Label $y_i = 1$ if x_i is an outlier, and $y_i = 0$ if it is not. The problem with this is that it is supervised. Hence, we need to know/define what outliers "look like", and new "types" of outliers may not be detected.

7. If we want to do linear regression with 1 feature, explain why it would or would not make sense to use gradient descent to compute the least squares solution.

Answer: It would not make sense to use gradient descent to compute least squares solution because with only 1 feature, it would take $O(n)$ time to compute solve least squares problem while with gradient descent it would take $O(nt)$ times

8. Why do we typically add a column of 1 values to X when we do linear regression? Should we do this if we're using decision trees?

Answer: Adding column of 1 values to X allows the least squares problem to compute a y-intercept, adjusting for biases. In decision trees, we don't make any speculation (adding bias) towards the target function, so we would not add a 1 column to X .

9. If a function is convex, what does that say about stationary points of the function? Does convexity imply that a stationary points exists?

Answer: If a function is convex, the stationary point of the function is a global minimizer. Convexity does not imply that a stationary point exists. For example, $f(x) = |x|$ is a convex function, but there is no x such that $f'(x) = 0$

10. Why do we need gradient descent for the robust regression problem, as opposed to just using the normal equations? Hint: it is NOT because of the non-differentiability. Recall that we used gradient descent even after smoothing away the non-differentiable part of the loss.

Answer: Gradient descent can be used to minimize the Huber loss in smooth approximations to the L1-norm and also fit a smooth approximation of the L_∞ norm.

11. What is the problem with having too small of a learning rate in gradient descent?

Answer: Small learning rate would mean the training process would take more time.

12. What is the problem with having too large of a learning rate in gradient descent?

Answer: Large learning rate could yield faster training process, but with sub-optimal results.

13. What is the purpose of the log-sum-exp function and how is this related to gradient descent?

Answer: The log-sum-exp function allows a smooth approximation of a non-smooth function. This allows gradient descent to find a global optimum. For example this could be used for the optimization on a loss-function using a max-norm.

14. What type of non-linear transform might be suitable if we had a periodic function?

Answer: We could include trigonometric transformation, since trigonometric functions are periodic.